




(a) All 2128 Channels Across 57 Layers



(b) GRA-L1 Agreement per Layer



(c) Selection Disagreement ($\uparrow$ Semantic / $\downarrow$ Redundant)



Statistical Summary ($p < 0.05$)

Total Channels Analyzed: 2128

Total Convolutional Layers: 57

GRA-L1 Correlation: $r = +0.075$

→ Weak positive correlation

→ Methods select DIFFERENT channels

Mean Agreement Rate: 51.9%

→ 48.1% of selections DIFFER

"Semantic Filters" Found: 530

(High GRA but Low L1)

"Redundant Channels" Found: 530

(High L1 but Low GRA)

CONCLUSION: GRA and L1 identify substantially different channels as important, supporting the hypothesis that weight magnitude does NOT equal semantic importance.